Introduction to Kubernetes

CODE

DEPLOY

BUILD

OPERATE

TEST

MONITOR

RELEASE

PLAN
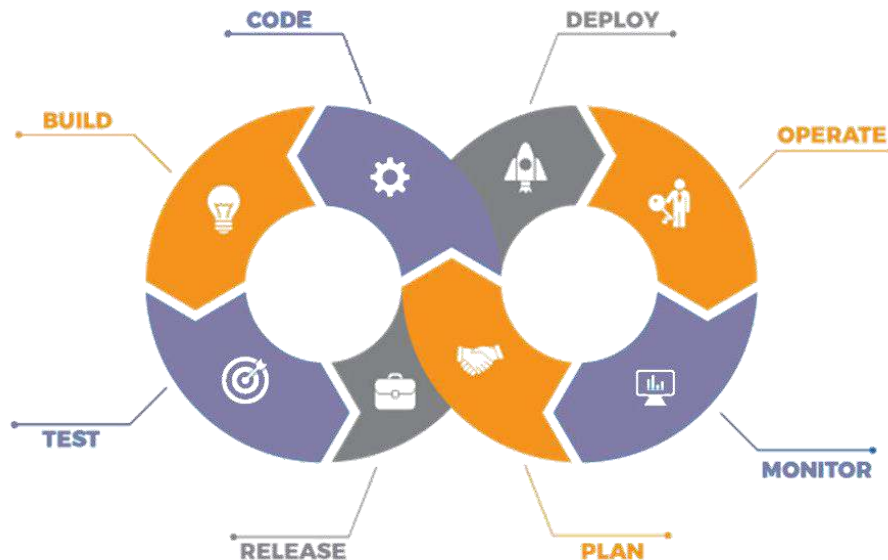
Agenda

01 INTRODUCTION TO KUBERNETES

02 DOCKER SWARM VS KUBERNETES

03 KUBERNETES ARCHITECTURE

04 KUBERNETES INSTALLATION

05 WORKING OF KUBERNETES

06 DEPLOYMENTS IN KUBERNETES

07 SERVICES IN KUBERNETES

08 INGRESS IN KUBERNETES

09 KUBERNETES DASHBOARD

Introduction to Kubernetes

Introduction to Kubernetes

⭐ Kubernetes is an open-source container orchestration software

⭐ It was originally developed by Google

⭐ It was first released on July 21st 2015

⭐ It is the ninth most active repository on GitHub in terms of number of commits

Features of Kubernetes

⭐ Pods	⭐ Service Discovery
⭐ Replication Controller	⭐ Networking
⭐ Storage Management	⭐ Secret Management
⭐ Resource Monitoring	⭐ Rolling Updates
⭐ Health Checks	

Docker Swarm vs Kubernetes

Docker Swarm vs Kubernetes

Source: trends.google.com

Docker Swarm vs Kubernetes

Docker Swarm

⭐ Easy to Install and Initialize

⭐ Faster when compared to Kubernetes
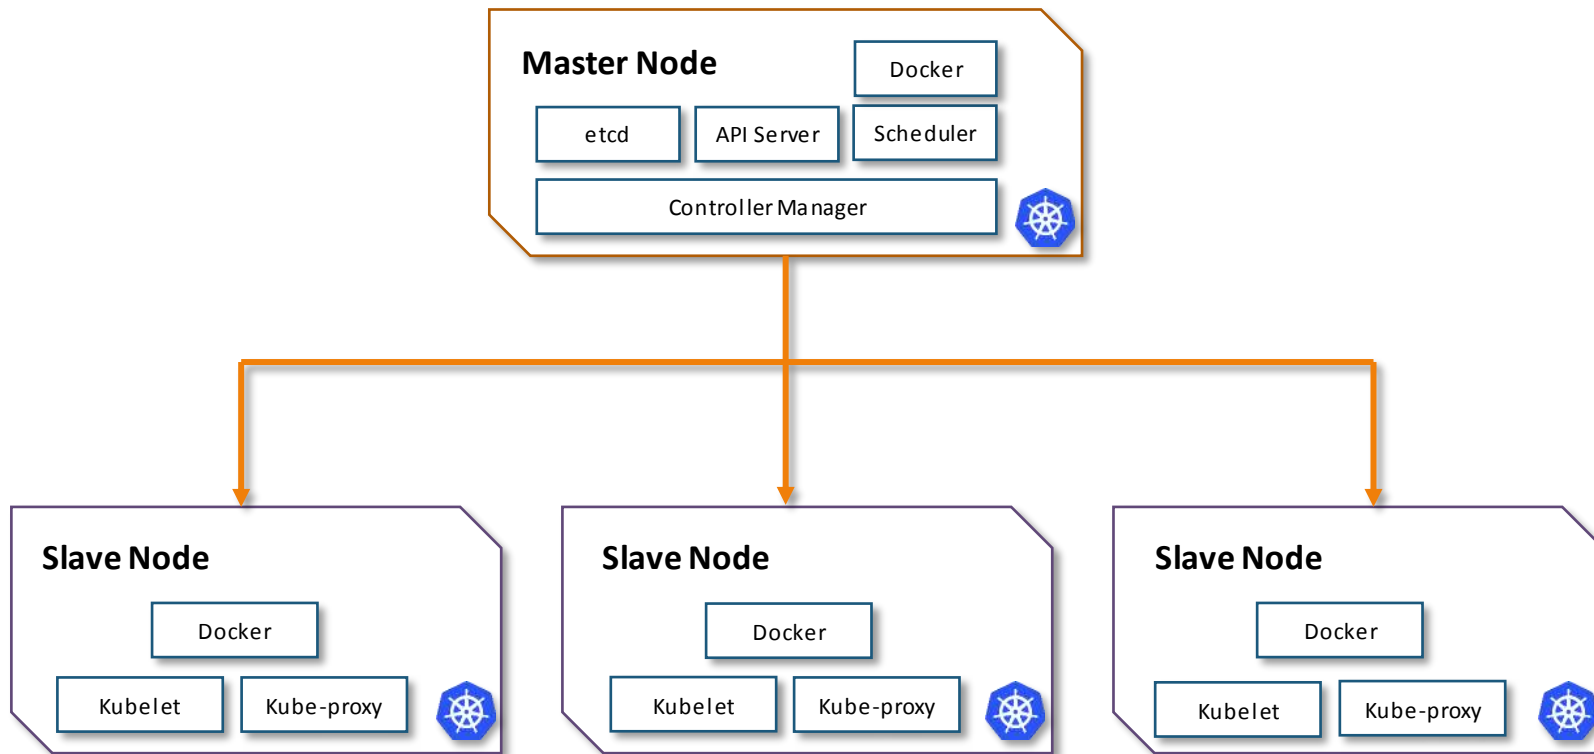
⭐ Not Reliable and has less features

Kubernetes

⭐ Complex Procedure to install Kubernetes

⭐ Slower when compared with Docker Swarm

⭐ More Reliable and comparatively has more features

Kubernetes Architecture

Kubernetes Architecture

Master Node

Slave Node

Slave Node

Slave Node

Kubernetes Architecture

Kubernetes Architecture – Master Components

Kubernetes Architecture – Master Components

IntelliPaat

etcd

API Server

Scheduler

Controller Manager

It is a highly available distributed key value store, which is used to store cluster wide secrets. It is only accessible by Kubernetes API server, as it has sensitive information.

Master Node

Docker

etcd

API Server

Scheduler

Controller Manager

Kubernetes Architecture – Master Components

etcd

API Server

Scheduler

Controller Manager

It exposes the Kubernetes API. The Kubernetes API is the front-end for Kubernetes Control Plane, and is used to deploy and execute all operations in Kubernetes

Master Node

Docker

etcd

API Server

Scheduler

Controller Manager

Kubernetes Architecture – Master Components

etcd

API Server

Scheduler

Controller Manager

The scheduler takes care of scheduling of all the processes, Dynamic Resource Management and manages present and future events on the cluster

Master Node

Docker

etcd

API Server

Scheduler

Controller Manager

Kubernetes Architecture – Master Components

etcd

API Server

Scheduler

Controller Manager

> The controller manager, runs all the controllers on the Kubernetes Cluster. Although each controller, is a separate process, but to reduce complexity, all the controllers are compiled into a single process. They are as follows:
> **Node Controller, Replication Controller, Endpoints Controller, Service Accounts and Token Controllers**

Master Node

Docker

etcd | API Server | Scheduler

Controller Manager

Kubernetes Architecture – Slave Components

Kubernetes Architecture – Master Components

Kubelet takes the specification from the API server, and ensures the application is running according to the specifications which were mentioned. Each node has it's kubelet service

Kubelet

Kube-Proxy

Slave Node

Docker

Kubelet Kube-proxy

Kubernetes Architecture – Master Components

This proxy service runs on each node and helps in making services available to the external host. It helps in connection forwarding to the correct resources, it is also capable of doing primitive load balancing

Kubelet

Kube-Proxy

Slave Node

Docker

Kubelet   **Kube-proxy**

Kubernetes Installation

Kubernetes Installation

There are numerous ways to install Kubernetes, following are some of the popular ways:

- **Kubeadm** – Bare Metal Installation

- **Minikube** – Virtualized Environment for Kubernetes

- **Kops** – Kubernetes on AWS

- **Kubernetes on GCP** – Kubernetes running on Google Cloud Platform

Hands-on: Installing Kubernetes using kubeadm

Working of Kubernetes

Working of Kubernetes

Pod – Replica 1

Pod – Replica 2

Pods can have one or more containers coupled together. They are the basic unit of Kubernetes. To increase High Availability, we always prefer pods to be in replicas

Pod – Replica 3

Working of Kubernetes

Services are used to load balance the traffic among the pods. It follows round robin distribution among the healthy pods
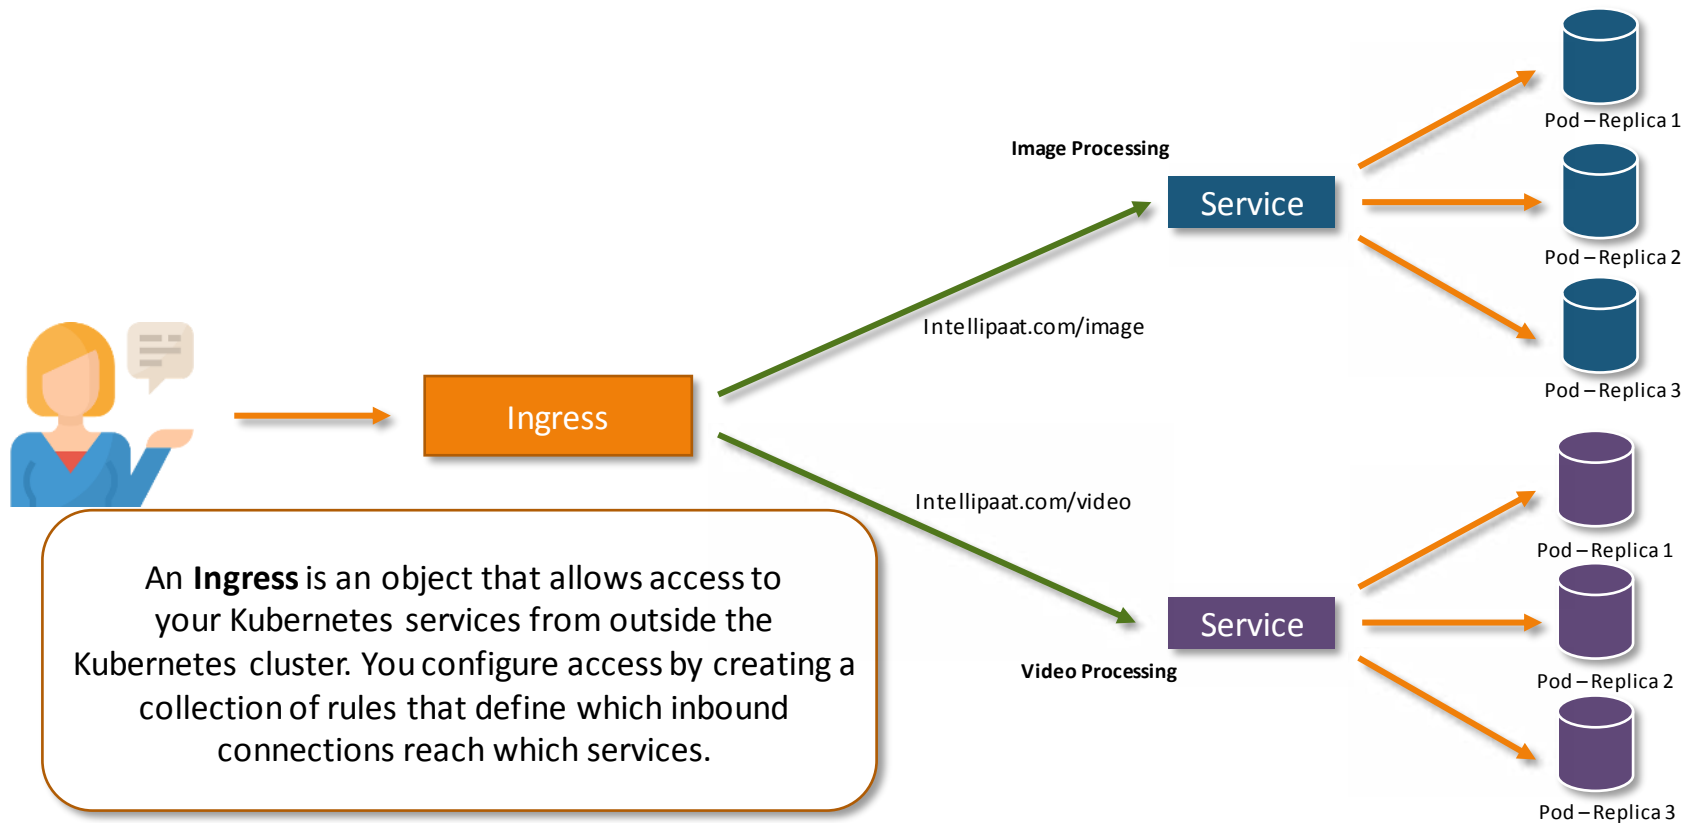
Service

Pod – Replica 1

Pod – Replica 2

Pod – Replica 3

Working of Kubernetes

Image Processing

Service

Pod – Replica 1

Pod – Replica 2

Pod – Replica 3

Intellipaat.com/image

Ingress

Intellipaat.com/video

An **Ingress** is an object that allows access to your Kubernetes services from outside the Kubernetes cluster. You configure access by creating a collection of rules that define which inbound connections reach which services.

Service

Video Processing

Pod – Replica 1

Pod – Replica 2

Pod – Replica 3

Deployments in Kubernetes

Deployments in Kubernetes

Deployment in Kubernetes is a controller which helps your applications reach the desired state, the desired state is defined inside the deployment file

Deployment

Pods

YAML Syntax for Deployments

This YAML file will deploy 3 pods for nginx, and maintain the desired state which is 3 pods, until this deployment is deleted

```yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80
```

Creating a Deployment

Once the file is created, to deploy this deployment use the following syntax:

```
ubuntu@ip-172-31-39-244: ~
ubuntu@ip-172-31-39-244:~$ kubectl create -f nginx.yaml
deployment.apps/nginx-deployment created
ubuntu@ip-172-31-39-244:~$
```

List the Pods

To view the pods, type the following command:

Syntax

kubectl get po

As you can see, the number of pods are matching with the number of replicas specified in the deployment file

Creating a Service

Creating a Service

A Service is basically a round-robin load balancer for all the pods, which match with it's name or selector. It constantly monitors the pods, in case a pod gets unhealthy, the service will start deploying the traffic to the other healthy pods.


Service

Pod – Replica 1

Pod – Replica 2

Pod – Replica 3

Service Types

ClusterIP: Exposes the service on cluster-internal IP

NodePort: Exposes the service on each Node's IP at a static port

LoadBalancer: Exposes the service externally using a cloud provider's load balancer.

ExternalName: Maps the service to the contents of the ExternalName

Pod – Replica 1

Service

Pod – Replica 2

Pod – Replica 3

Creating a NodePort Service

We can create a NodePort service using the following syntax:

Syntax

```
kubectl create service nodeport <name-of-service> --tcp=<port-of-service>:<port-of-container>
```

```
ubuntu@ip-172-31-39-244: ~
ubuntu@ip-172-31-39-244:~$ kubectl create service nodeport nginx --tcp=80:80
service/nginx created
ubuntu@ip-172-31-39-244:~$
```

Creating a NodePort Service

To know the port, on which the service is being exposed type the following command:

> **Syntax**
>
> kubectl get svc nginx

```
ubuntu@ip-172-31-39-244: ~

ubuntu@ip-172-31-39-244:~$ kubectl get svc nginx
NAME      TYPE        CLUSTER-IP       EXTERNAL-IP    PORT(S)         AGE
nginx     NodePort    10.103.235.81    <none>         80:32043/TCP    114s
ubuntu@ip-172-31-39-244:~$
```
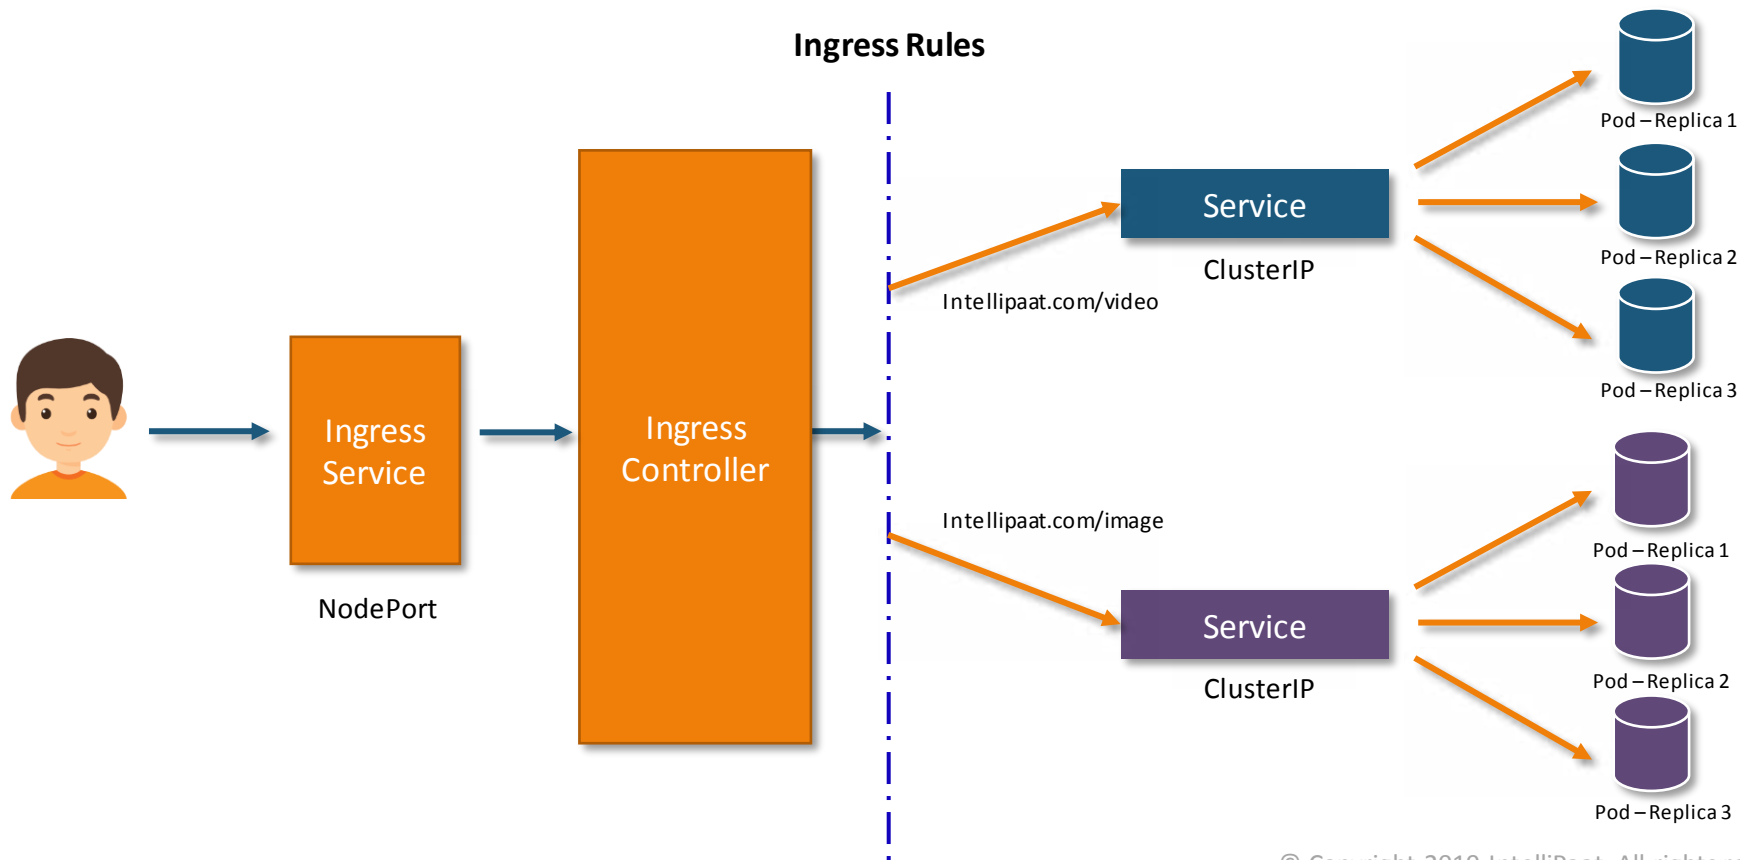
Creating a NodePort Service

To know the port, on which the service is being exposed type the following command:

Syntax

kubectl get svc nginx

```
ubuntu@ip-172-31-39-244: ~
ubuntu@ip-172-31-39-244:~$ kubectl get svc nginx
NAME     TYPE       CLUSTER-IP      EXTERNAL-IP    PORT(S)        AGE
nginx    NodePort   10.103.235.81   <none>         80:32043/TCP   114s
ubuntu@ip-172-31-39-244:~$
```

Creating an Ingress

What is an Ingress?

Kubernetes ingress is a collection of routing rules that govern how external users access services running in a Kubernetes cluster.

Intellipaat.com/image

Service

Ingress

Intellipaat.com/video

Service

What is an Ingress?

Ingress Rules

Intellipaat.com/video

Intellipaat.com/image

Service — ClusterIP

Service — ClusterIP

Ingress Service — NodePort

Ingress Controller

Pod − Replica 1
Pod − Replica 2
Pod − Replica 3

Pod − Replica 1
Pod − Replica 2
Pod − Replica 3

Installing Ingress Controller

We will be using the nginx ingress controller, for our demo. We can download it from the following link:

Link

https://github.com/kubernetes/ingress-nginx/blob/master/docs/deploy/index.md

Define Ingress Rules

The following rule, will redirect traffic which asks for /foo to nginx service. All the other requests, will be redirected to ingress controller's default page

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: simple-fanout-example
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
     paths:
     - path: /foo
       backend:
         serviceName: nginx
         servicePort: 80
```

Deploying Ingress Rules

To deploy the ingress rules, we use the following syntax:

kubectl create –f ingress.yaml

```
ubuntu@ip-172-31-17-194: ~
ubuntu@ip-172-31-17-194:~$ kubectl create -f ingress.yaml
ingress.extensions/simple-fanout-example created
ubuntu@ip-172-31-17-194:~$
```

Viewing Ingress Rules

To deploy the ingress rules, we use the following syntax:

Syntax

kubectl get ing

```
ubuntu@ip-172-31-17-194: ~
ubuntu@ip-172-31-17-194:~$ kubectl get ing
NAME                      HOSTS    ADDRESS    PORTS    AGE
simple-fanout-example     *                   80       2m5s
ubuntu@ip-172-31-17-194:~$
```

Kubernetes Dashboard

Kubernetes Dashboard

Dashboard is a web-based Kubernetes user interface. You can use Dashboard to deploy containerized applications to a Kubernetes cluster, troubleshoot your containerized application, and manage the cluster resources.
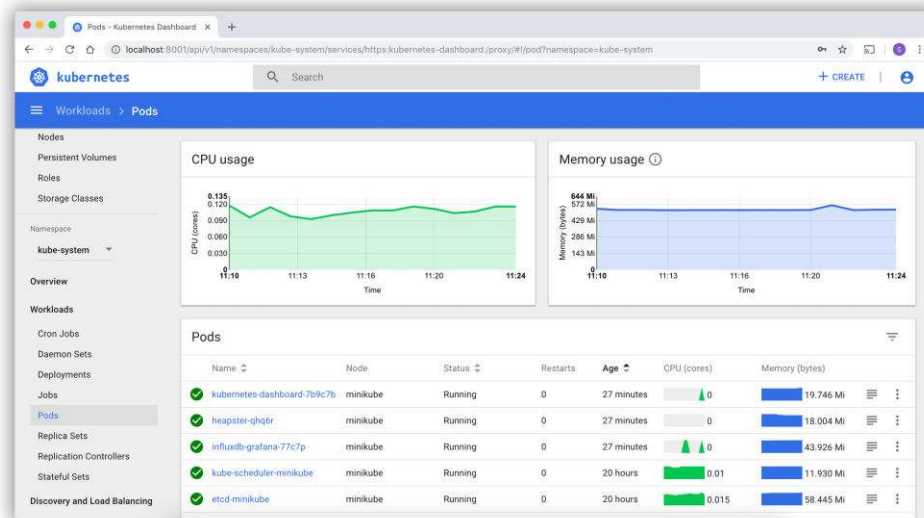
Installing Kubernetes Dashboard

To install Kubernetes Dashboard, execute the following command:

kubectl create -f
https://raw.githubusercontent.com/kubernetes/dashboard/master/aio/deploy/recommended/kubernetes-dashboard.yaml

Accessing Kubernetes Dashboard

Change the service type for Kubernetes-Dashboard to Nodeport

Syntax

kubectl -n kube-system edit service kubernetes-dashboard

Logging into Kubernetes Dashboard

1. Check the NodePort from the kubernetes-dashboard service
2. Browse to your cluster on the internet browser, and enter the IP address
3. Click on Token, it will ask you for the token entry
4. Generate a token using the following command

```
$ kubectl create serviceaccount cluster-admin-dashboard-sa
$ kubectl create clusterrolebinding cluster-admin-dashboard-sa \
  --clusterrole=cluster-admin \
  --serviceaccount=default:cluster-admin-dashboard-sa

$ TOKEN=$(kubectl describe secret $(kubectl -n kube-system get secret | awk '/^cluster-admin-dashboard-sa-token-/{print $1}') | awk '$1=="token:"{print $2}')

$ echo $TOKEN
```

5. Finally, enter the token and login to your dashboard

India : +91-7847955955

US : 1-800-216-8930 (TOLL FREE)

support@intellipaat.com

24X7 Chat with our Course Advisor